

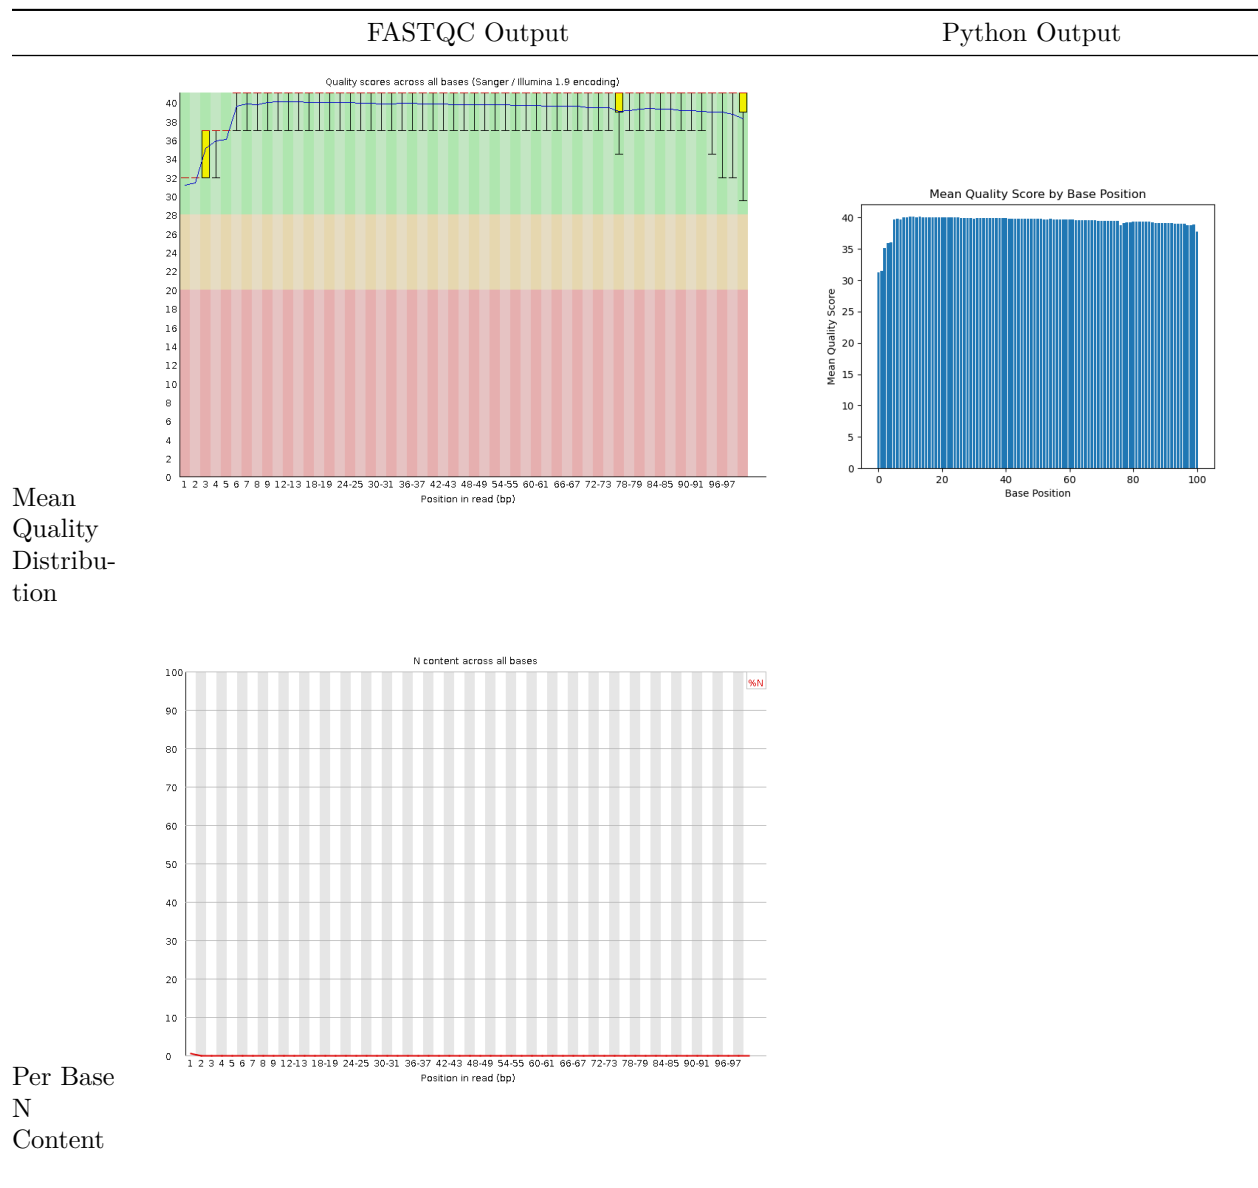
QAA_Sam_Kupp

2022-09-07

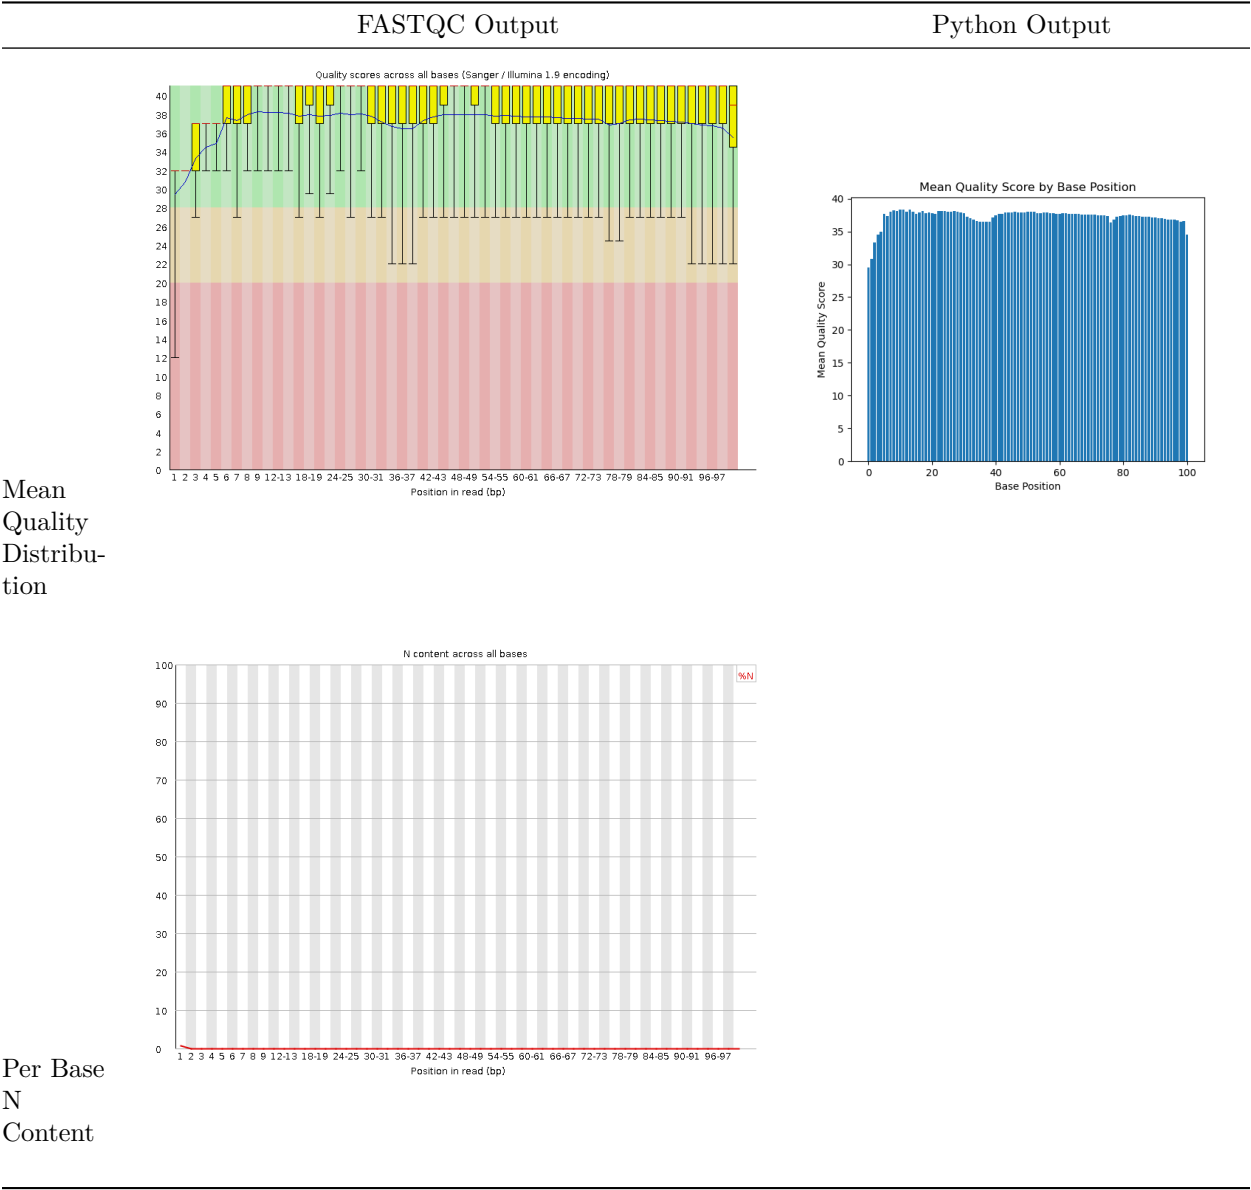
Part 1- Read Quality Score Distributions

Quality Score Distributions from FASTQC

23_4A_control_S17_L008 Read 1



23_4A_control_S17_L008 Read 2





31_4F_fox_S22_L008 Read 1


FASTQC Output

Python Output

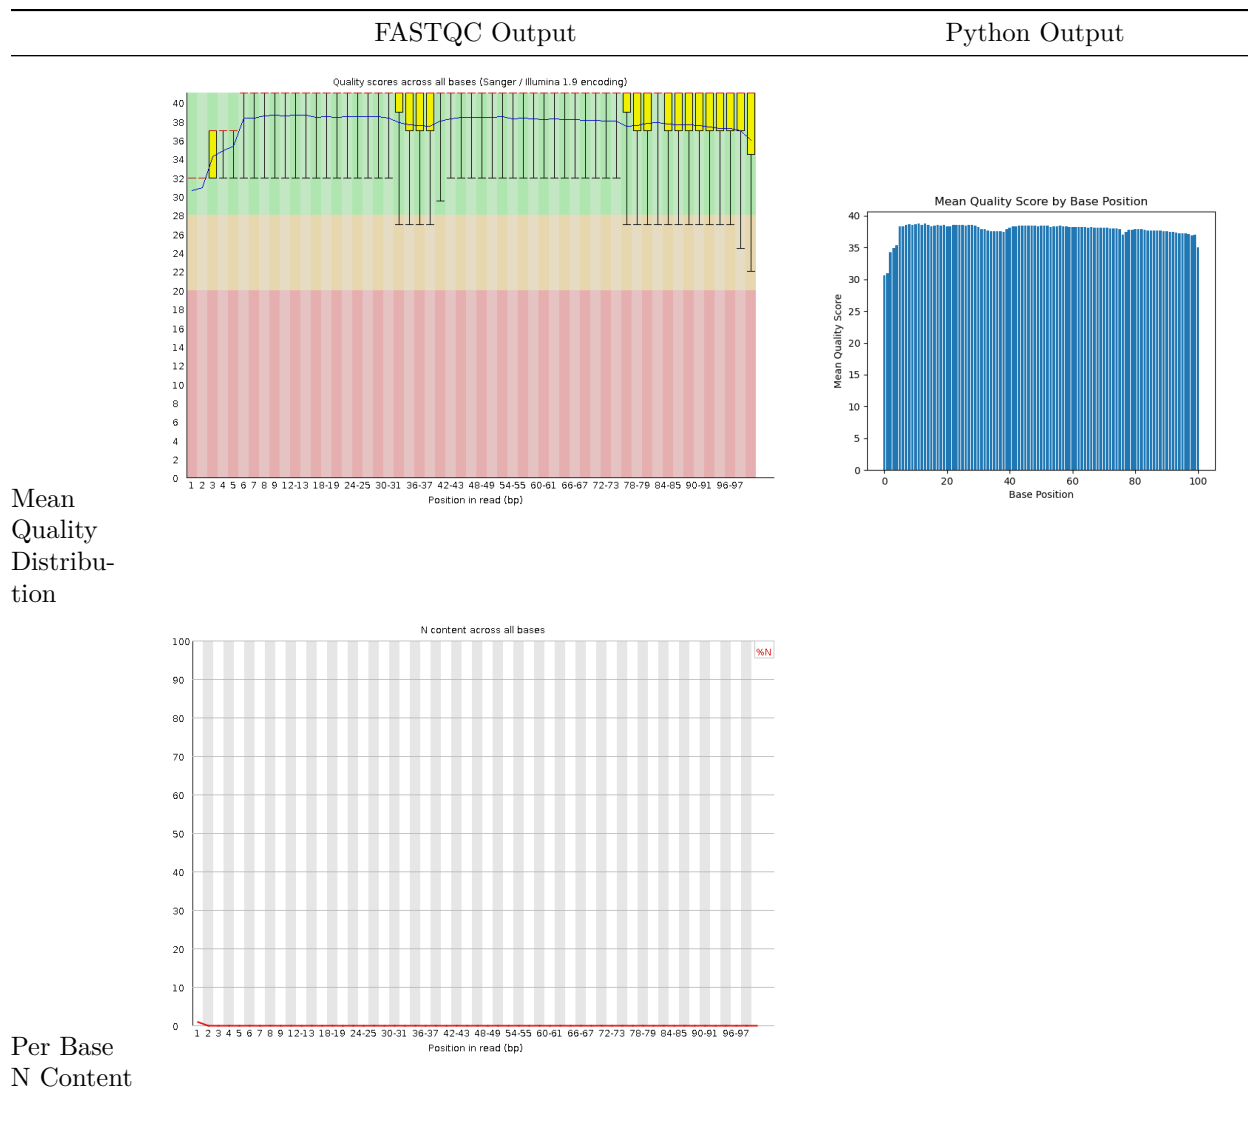
Mean
Quality
Distribu-
tion



Per Base
N Content



31_4F_fox_S22_L008 Read 2



The per base n content plots are consistent with the per base quality plots. In each of the 4 files, the only base position with a significant proportion of reads that are ‘N’, is the first. This is also the position with the lowest average quality.

As expected, the score distributions from FASTQC and the Python script are identical, with any visual discrepancies likely due to the plotting geometry of bar vs line graph.

Run Times for Quality Distribution Plots

| FASTQ File | Python Script Run Time (s) | FASTQC Run Time (s) |
|-------------------------------|----------------------------|---------------------|
| 23_4A_control_S17_L008 Read 1 | 1376.91 | 166.47 |
| 23_4A_control_S17_L008 Read 2 | 1408.47 | 166.57 |
| 31_4F_fox_S22_L008 Read 1 | 115.9 | 16.56 |
| 31_4F_fox_S22_L008 Read 2 | 117.31 | 17.21 |

As can be seen from the table, FASTQC is much faster than my Python Script. This is likely at least in

part because Python is an extremely flexible, but slow language. It is also possible that unzipping protocol that FASTQC uses is more efficient than gzip in Python.

The library qualities are as generally expected. Each of the first reads is relatively low quality, the later ~80 reads are between 37-40 quality, but taper towards the later reads in the file. Additionally, the reverse reads in each file set are lower quality as the DNA has likely degraded before these reads were sequenced.

Part 2- Adapter trimming comparison

Cut Adapt

CutAdapt Percent of Reads Trimmed

| FASTQ File Pair | Percent of Filtered Base Pairs |
|------------------------|--------------------------------|
| 23_4A_control_S17_L008 | 96.8 |
| 31_4F_fox_S22_L008 | 99.7 |

Adapter 1 Check

```
zcat 23_4A_control_S17_L008_R1_001.fastq.gz | head -100000 |grep -c "AGATCGGAAGAGCACACGTCTGAACTCCAGTCA"
29
zcat 23_4A_control_S17_L008_R2_001.fastq.gz | head -100000 |grep -c "AGATCGGAAGAGCACACGTCTGAACTCCAGTCA"
0
```

```
zcat 31_4F_fox_S22_L008_R1_001.fastq.gz | head -100000|grep -c "AGATCGGAAGAGCACACGTCTGAACTCCAGTCA"
629
zcat 31_4F_fox_S22_L008_R2_001.fastq.gz | head -100000|grep -c "AGATCGGAAGAGCACACGTCTGAACTCCAGTCA"
0
```

Adapter 2 Check

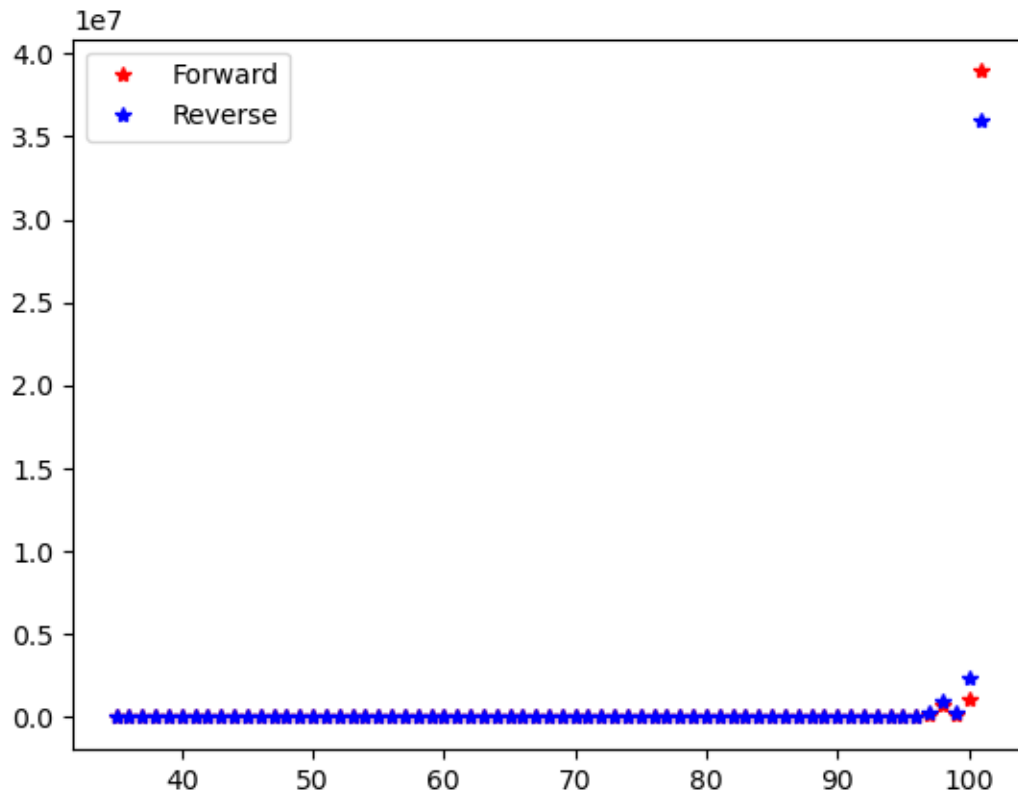
```
zcat 23_4A_control_S17_L008_R1_001.fastq.gz | head -100000 |grep -c "AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT"
0
zcat 23_4A_control_S17_L008_R2_001.fastq.gz | head -100000 |grep -c "AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT"
29
```

```
zcat 31_4F_fox_S22_L008_R1_001.fastq.gz | head -100000| grep -c "AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT"
0
zcat 31_4F_fox_S22_L008_R2_001.fastq.gz | head -100000| grep -c "AGATCGGAAGAGCGTCGTGTAGGGAAAGAGTGT"
630
```

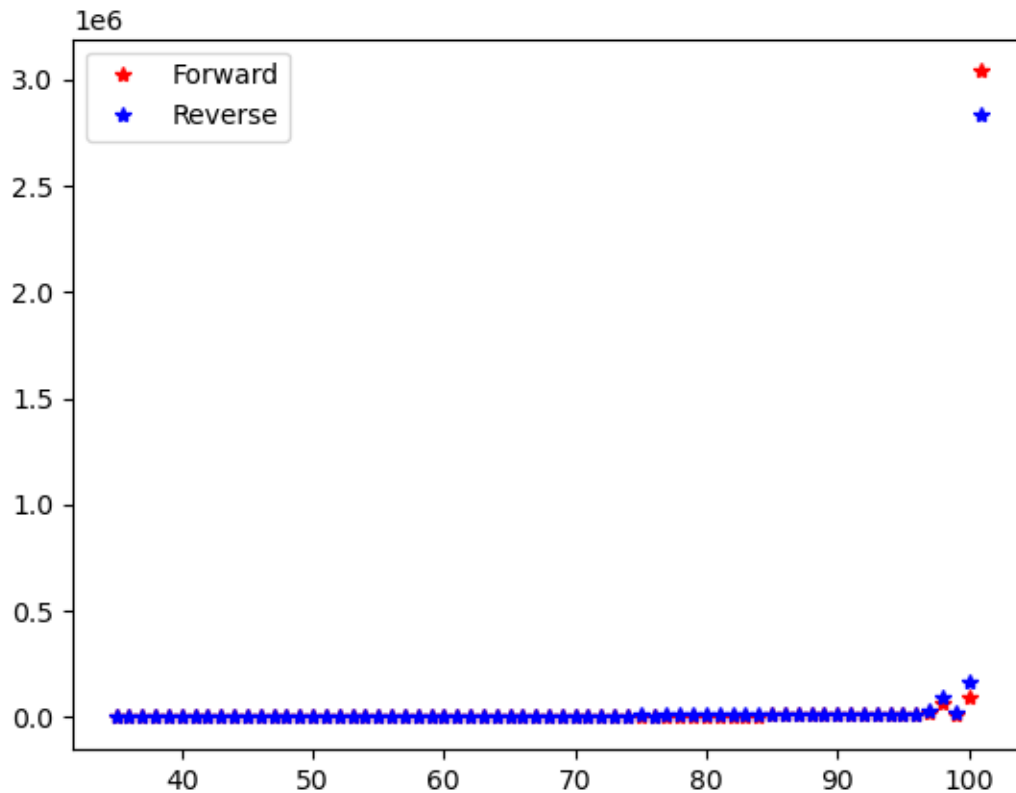
The above commands were used to search for each adapter sequence in each of the read files. Here the first 25000 sequences were examined. As expected adapter sequence 1 was found only in read 1 and adapter sequence 2 was found only in read 2. This confirms that the reads were properly oriented.

Trimmomatic

23_4A_control_S17_L008 Trimmed Sequence Lengths



31_4F_fox_S22_L008 Trimmed Sequence Lengths



As expected, the forward reads exhibit more reads at 101 base pairs than the reverse read. This is likely due to degradation of the DNA on the sequencer and the quality trimming done during Trimmomatic. Both distributions are heavily peaked at 101 base pairs, indicating that both reads are high quality.

Part 3- Alignment and strand-specificity

After building a STAR database and aligning the trimmed reads to it, the results are as shown below:

Mapped Reads from Python Script

| Read File | Number of Reads Mapped | Number of Reads Not Mapped |
|------------------------|------------------------|----------------------------|
| 23_4A_control_S17_L008 | 79473028 | 4640124 |
| 31_4F_fox_S22_L008 | 6969880 | 225936 |

HTSeq Counts Output

Mapped Reads from HTSeq

| Count File | Number of Reads Mapped | Total Number of Reads | Percentage of Reads Mapped |
|------------------------|---------------------------|--------------------------|-------------------------------|
| 23_4A_control_S17_L008 | 65634193 | 81286685 | 80.7441 |
| stranded=reverse | | | |
| 23_4A_control_S17_L008 | 2261967 | 81286685 | 2.7827 |
| stranded=yes | | | |
| 31_4F_fox_S22_L008 | 5691153 | 6695862 | 84.9951 |
| stranded=reverse | | | |
| 31_4F_fox_S22_L008 | 291122 | 6695862 | 4.34779 |
| stranded=yes | | | |

Reads Mapped, Total Number of Reads and Percentage of Reads Mapped Calculated with:

```
for file in `ls ~/Documents/BI_622/QAA/*.tsv`; do echo $file ; awk '{sum1 += $2}\
$1~"ENSMU"{sum2 += $2} END {print sum2 "\t" sum1 "\t" sum2/sum1*100}' $file ; done
```

```
## /Users/sam/Documents/BI_622/QAA/23_4A_control_S17_L008.reverse.count.tsv
## 65634193 81286685 80.7441
## /Users/sam/Documents/BI_622/QAA/23_4A_control_S17_L008.stranded.count.tsv
## 2261967 81286685 2.7827
## /Users/sam/Documents/BI_622/QAA/31_4F_fox_S22_L008.reverse.count.tsv
## 5691153 6695862 84.9951
## /Users/sam/Documents/BI_622/QAA/31_4F_fox_S22_L008.stranded.count.tsv
## 291122 6695862 4.34779
```

From the percentage of reads mapped, we can determine that the library is strand specific. Because roughly 80% of reads are mapped when stranded=reverse in the HTSeq command but only less than 5% are mapped when stranded=yes. This means that a specific strand was sequenced, not both as in an unstranded library prep.